

I. Overview	1
II. Implement	1
1. Remote Backend configuration	1
2. Prometheus & Grafana addition	1
3. Execution	5
4. Verify deployed resource	6

Deploy AKS Using Terraform

I. Overview

- The goal of this document is to demonstrate how to deploy AKS, VNet, Prometheus, and Grafana to Azure using Terraform.
- In this demo ([my repository](#)), I leveraged [the original provided repository](#) and added some new resources to meet the assignment requirements.

II. Implement

1. Remote Backend configuration

In the top of **./main.tf** add following block

```
terraform {  
  backend "azurerm" {  
    resource_group_name = "myTFResourceGroup"  
    storage_account_name = "minhpraticekubernetes"  
    container_name       = "tfstate"  
    key                  = "terraform.tfstate"  
  }  
}
```

This block configures the **Terraform backend** to use **Azure Storage** for remote state management.

It stores the Terraform state file (terraform.tfstate) in the specified storage

account, container, and resource group, ensuring state consistency and team collaboration.

2. ACR

In the end of **.modules/aks/main.tf** add a block to define ACR

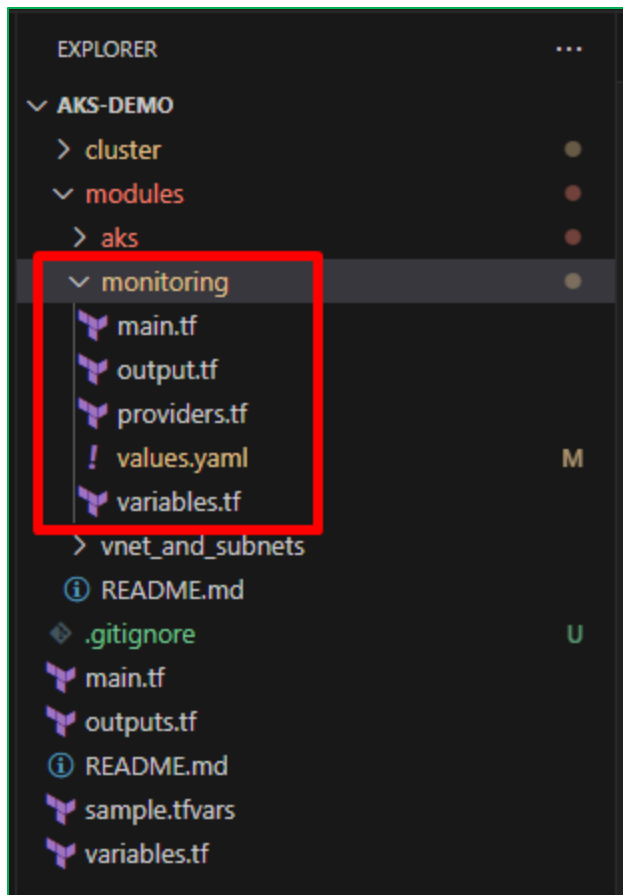
```
resource "azurerm_container_registry" "minhtestacr" {
  name                = "minhtestacr"
  resource_group_name = var.resource_group_name
  location            = var.region
  sku                 = "Standard"
}

resource "azurerm_role_assignment" "minh_aks_acr_pull" {
  principal_id =
azurerm_kubernetes_cluster.k8s.kubelet_identity[0].object_id
  role_definition_name = "AcrPull"
  scope              =
azurerm_container_registry.minhtestacr.id
  skip_service_principal_aad_check = true
}
```

block `azurerm_role_assignment` will create a IAM role to allow AKS pull image from ACR

3. Prometheus & Grafana addition

Add module monitoring which define code to deploy Prometheus and Grafana using helm



```
# main.tf
resource "helm_release" "kube_prometheus_stack" {
  # Helm repository details
  name       = "kube-prometheus-stack"
  repository = "https://prometheus-community.github.io/helm-charts"
  chart      = "kube-prometheus-stack"
  version    = "59.0.0" # Use a stable, recent version
  namespace  = "monitoring"
  create_namespace = true

  values = [
    templatefile("${path.module}/values.yaml", {
      ip = ""
    })
  ]

  # Optional: Wait for all resources to be ready (can take a few minutes)
  timeout = 600
  wait    = true
}
```

This Helm release installs the **Kube Prometheus Stack** chart into the monitoring namespace on AKS. It deploys Prometheus, Grafana, and related monitoring components using a specified Helm chart version from the Prometheus Community repository.

```
# providers.tf
# Configure the Kubernetes Provider
provider "kubernetes" {
  host                = var.host
  client_certificate   = base64decode(var.client_certificate)
  client_key           = base64decode(var.client_key)
  cluster_ca_certificate = base64decode(var.cluster_ca_certificate)
}

# Configure the Helm Provider
provider "helm" {
  kubernetes = {
    host                = var.host
    client_certificate   = base64decode(var.client_certificate)
    client_key           = base64decode(var.client_key)
    cluster_ca_certificate = base64decode(var.cluster_ca_certificate)
  }
}
```

This configuration defines the **Kubernetes** and **Helm** providers used to interact with the AKS cluster.

It connects to the cluster using credentials (host, client certificate, client key, and CA certificate) passed through Terraform variables.

```
# values.yaml
prometheus:
  prometheusSpec:
    podMonitorSelectorNilUsesHelmValues: false
    serviceMonitorSelectorNilUsesHelmValues: false
    ruleSelectorNilUsesHelmValues: false

grafana:
  adminUser: admin
  adminPassword: admin
  grafana.ini:
```

```
server:
  root_url: "%(protocol)s://%(domain)s/grafana/"
```

This **values.yaml** snippet customizes the **Kube Prometheus Stack** deployment.

- Under `prometheus.prometheusSpec`, it ensures Prometheus recognizes **PodMonitors**, **ServiceMonitors**, and **Rules** created outside of the Helm release (by disabling the default Helm-only selectors).
- The `grafana` section sets up Grafana with default admin credentials and defines a custom `root_url` for accessing the Grafana UI (then we can access Grafana via <Ingress controller IP>/grafana)

```
# variables.tf
variable "host" {
  type      = string
  description = "aks cluster host"
}

variable "client_certificate" {
  type      = string
  description = "aks cluster client certificate"
}

variable "client_key" {
  type      = string
  description = "aks cluster client key"
}

variable "cluster_ca_certificate" {
  type      = string
  description = "aks cluster cluster ca certificate"
}
```

This file defines the Terraform input variables required to authenticate and connect to the AKS cluster.

In **./modules/aks/outputs.tf** add some outputs:

```
output "client_certificate" {
  description = "Base64 encoded public certificate used by clients to authenticate to the Kubernetes cluster."
  value      =
    azurerm_kubernetes_cluster.k8s.kube_config.0.client_certificate
```

```

    sensitive = true
  }
  output "host" {
    description = "Kubernetes cluster host."
    value      = azurerm_kubernetes_cluster.k8s.kube_config.0.host
    sensitive  = true
  }

  output "client_key" {
    description = "Base64 encoded public client_key."
    value      = azurerm_kubernetes_cluster.k8s.kube_config.0.client_key
    sensitive  = true
  }

  output "cluster_ca_certificate" {
    description = "Base64 encoded public cluster_ca_certificate."
    value      =
    azurerm_kubernetes_cluster.k8s.kube_config.0.cluster_ca_certificate
    sensitive  = true
  }

```

Finally, in **./cluster/main.tf** add section to call the monitoring module and pass value for variables

```

module "monitoring" {
  # invoke monitoring module under modules directory
  source = "../modules/monitoring"

  host = module.aks_with_node_group.host
  client_certificate=module.aks_with_node_group.client_certificate
  client_key=module.aks_with_node_group.client_key
  cluster_ca_certificate=module.aks_with_node_group.cluster_ca_certificate
}

```

4. Execution

Cd to root of repository folder, execute

terraform init

terraform apply -var-file="sample.tfvars"

and then enter “yes” to confirm start process

```
PS /home/minh> cd ./aks-demo/
PS /home/minh/aks-demo> terraform init
Initializing the backend...
Initializing modules...
Initializing provider plugins...
- Reusing previous version of hashicorp/helm from the dependency lock file
- Reusing previous version of hashicorp/kubernetes from the dependency lock file
- Reusing previous version of hashicorp/azurerm from the dependency lock file
- Using previously-installed hashicorp/helm v3.1.0
- Using previously-installed hashicorp/kubernetes v2.38.0
- Using previously-installed hashicorp/azurerm v3.107.0

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
PS /home/minh/aks-demo> terraform apply -var-file="sample.tfvars"
Acquiring state lock. This may take a few moments...
module.cluster.module.vnet_with_subnets.azure_rm_virtual_network.minh_az_vnet: Refresh
```

5. Verify deployed resource

Use az cli command to access AKS

az aks get-credentials --resource-group "your_resource_group" --name "your_AKS"

```
Switch to Bash Restart Manage files New session Editor Web preview Settings Help
PS /home/minh> az aks get-credentials --resource-group "myTFResourceGroup" --name "minh-test-aks"
Merged "minh-test-aks" as current context in /home/minh/.kube/config
PS /home/minh> 
```

Use **kubectl get all -n monitoring** to verify Prometheus and Grafana

```
PS /home/minh> kubectl get all -n monitoring
```

NAME	READY	STATUS	RESTARTS	AGE
pod/alertmanager-kube-prometheus-stack-alertmanager-0	2/2	Running	0	4h29m
pod/kube-prometheus-stack-grafana-7d47bc94b9-g9q2s	3/3	Running	0	3h18m
pod/kube-prometheus-stack-kube-state-metrics-5b8c8b74d5-rgxxl	1/1	Running	0	4h29m
pod/kube-prometheus-stack-operator-768986f7c5-881jb	1/1	Running	0	4h29m
pod/kube-prometheus-stack-prometheus-node-exporter-72rjz	1/1	Running	0	4h29m
pod/kube-prometheus-stack-prometheus-node-exporter-ptxz2	1/1	Running	0	4h29m
pod/prometheus-kube-prometheus-stack-prometheus-0	2/2	Running	0	4h29m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/alertmanager-operated	ClusterIP	None	<none>	9093/TCP, 9094/TCP, 9094/UDP	4h29m
service/kube-prometheus-stack-alertmanager	ClusterIP	10.0.74.156	<none>	9093/TCP, 8080/TCP	4h29m
service/kube-prometheus-stack-grafana	ClusterIP	10.0.161.193	<none>	80/TCP	4h29m
service/kube-prometheus-stack-kube-state-metrics	ClusterIP	10.0.151.239	<none>	8080/TCP	4h29m
service/kube-prometheus-stack-operator	ClusterIP	10.0.39.166	<none>	443/TCP	4h29m
service/kube-prometheus-stack-prometheus	ClusterIP	10.0.146.192	<none>	9090/TCP, 8080/TCP	4h29m
service/kube-prometheus-stack-prometheus-node-exporter	ClusterIP	10.0.204.239	<none>	9100/TCP	4h29m
service/prometheus-operated	ClusterIP	None	<none>	9090/TCP	4h29m

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR
daemonset.apps/kube-prometheus-stack-prometheus-node-exporter	2	2	2	2	2	kubernetes.io/os=linux

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/kube-prometheus-stack-grafana	1/1	1	1	4h29m
deployment.apps/kube-prometheus-stack-kube-state-metrics	1/1	1	1	4h29m
deployment.apps/kube-prometheus-stack-operator	1/1	1	1	4h29m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/kube-prometheus-stack-grafana-544c99459	0	0	0	4h29m
replicaset.apps/kube-prometheus-stack-grafana-7d47bc94b9	1	1	1	3h18m
replicaset.apps/kube-prometheus-stack-kube-state-metrics-5b8c8b74d5	1	1	1	4h29m
replicaset.apps/kube-prometheus-stack-operator-768986f7c5	1	1	1	4h29m